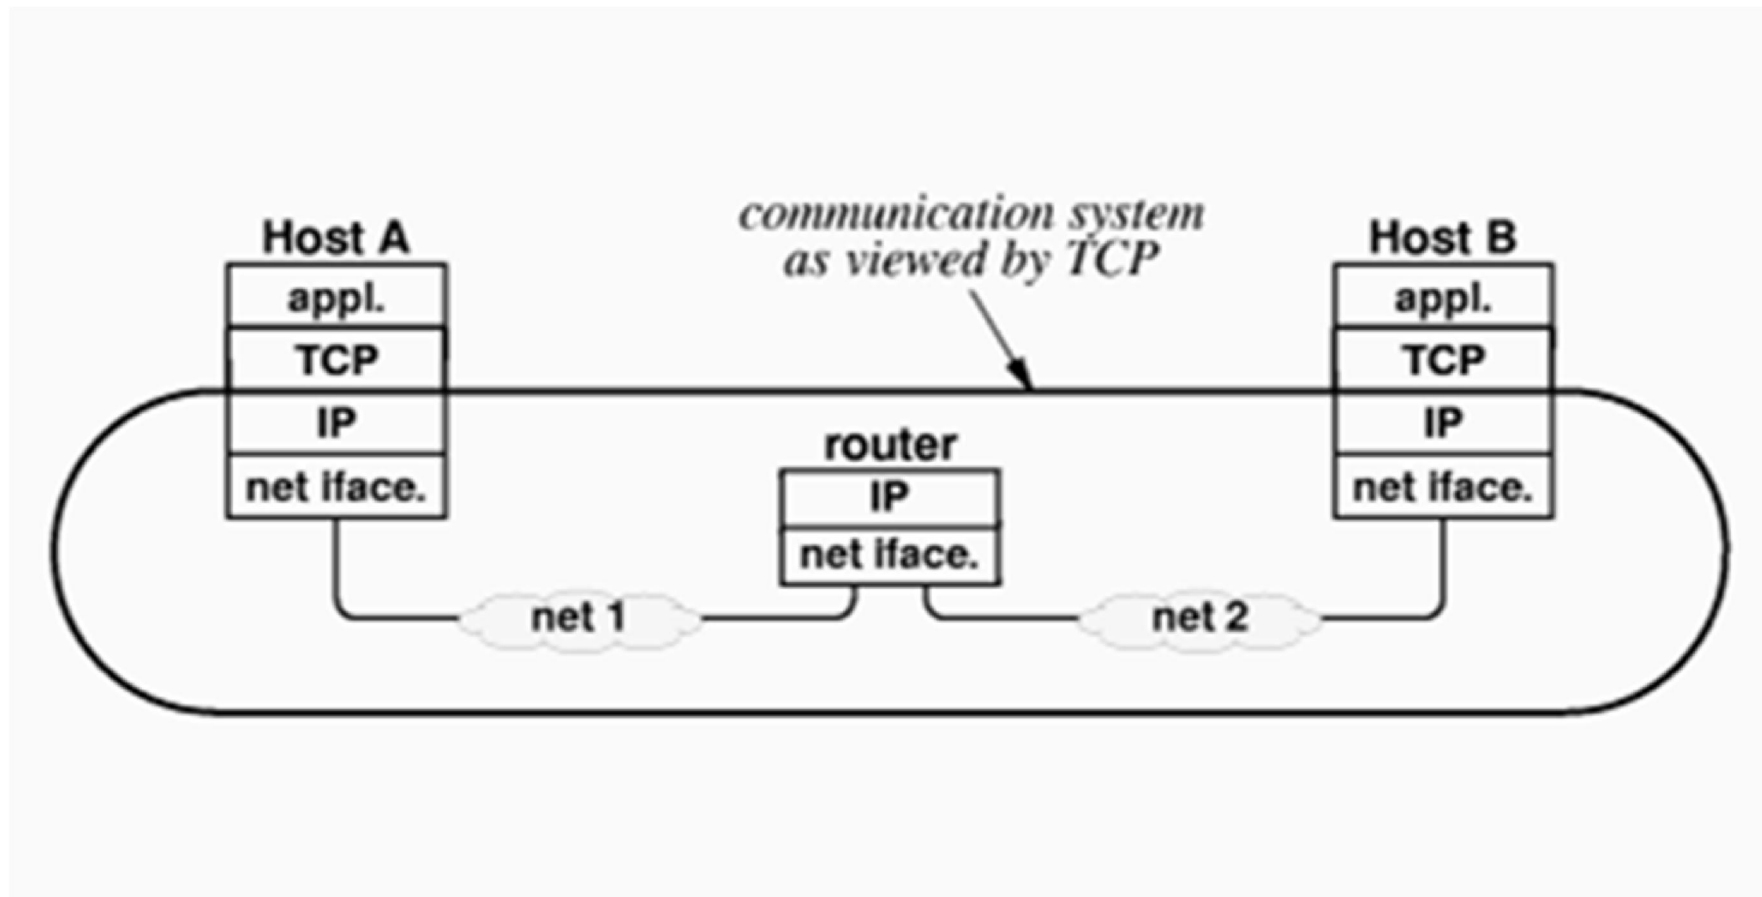


The Transport Layer

- Having examined the TCP/IP Protocol Reference the focus now shifts to the TCP layer.
 - This layer is the heart of the whole protocol hierarchy.
- It sits below the Application Layer providing a ‘Transport’ service to the applications that wish to communicate across a Network.

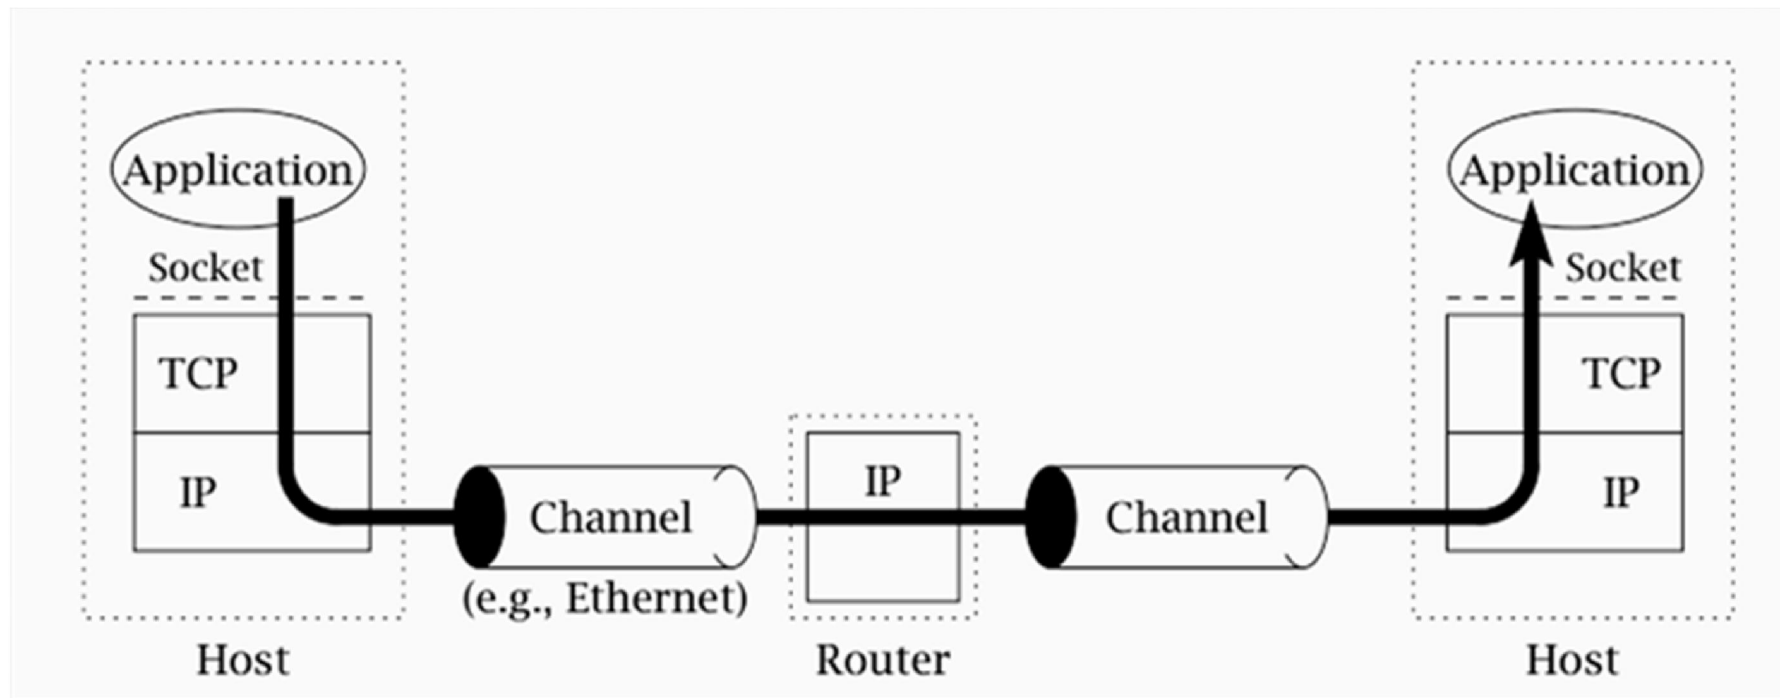
The Transport Layer in Context



The Transport Layer

- TCP provides a reliable, cost-effective end-to-end data transport service *independently* of the physical network(s).
- It is important to realise that a *service* is very different to a *protocol*, although they are frequently confused.
- The next slide shows how the Transport Layer views the lower network interface layers.

The Transport Layer's view of the Network



Services

- Each layer of the reference model provides a set of functionality to the layer immediately above:
 - This set of functionality is known as **The Services**.
 - The layer below is known as the **Service Provider** and the layer above is known as the **Service User**.
- The Services are accessed through the interface between the layers.

Protocols

- **Protocols** on the other hand are how the Services are implemented.
- Typically a Protocol specifies a *framing structure* which will include a number of fields containing Control data:
 - Generically this framing structure is known as a Protocol Data Unit (**PDU**)
 - Examples include Data Link Frame, IP Datagram etc.

Protocols

- The **Protocol** will also typically specify a procedure for interpreting and responding to the Control data within the PDU:
 - For instance, if a service provides for reliable transfer of data then there must be some means specified within the protocol for tracking and recovering from data loss.
 - In this instance part of the PDU Control field will include numbering e.g. byte numbers, frame numbers etc.

Protocols

- The Protocol will also specify how data in the Control Fields are to be interpreted and responded to if necessary.
 - e.g. for missing frame return a REJ message.

The TCP Transport Service Offering

- The TCP Transport Service has the following characteristics:
 - **Connection Orientation:** Before two applications entities can communicate they must first establish a connection.
 - **Point-To-Point Communication:** Each TCP connection has exactly two endpoints.
 - **Complete Reliability:** TCP guarantees that the data will be delivered exactly as sent i.e. no data missing or out of sequence
 - **Full Duplex Communication:** A TCP connection allows data to flow in either direction
 - TCP buffers outgoing and incoming data
 - This allows applications to continue executing other code whilst the data is being transferred

The TCP Transport Service Offering

- ***Stream Interface***: The source application sends a *continuous* sequence of octets across a connection
 - The data is passed *en bloc* to TCP for delivery
 - TCP does not guarantee to deliver the data in the same size pieces that it was transferred by the source application.
- ***Reliable Connection Startup***: TCP both applications to agree to any new connection
- ***Graceful Connection Shutdown***: Either can request a connection to be shut down
 - TCP guarantees to deliver all the data reliably before closing the connection

The Transport Service

- The TCP transport service is offered to a *user process* that exists within the application layer:
 - This *user process* is considered a ***Transport Service User***,
 - The service is typically offered through a set of **primitives** across the interface between the layers,
 - Calls to these primitives cause the transport service provider to perform some action.

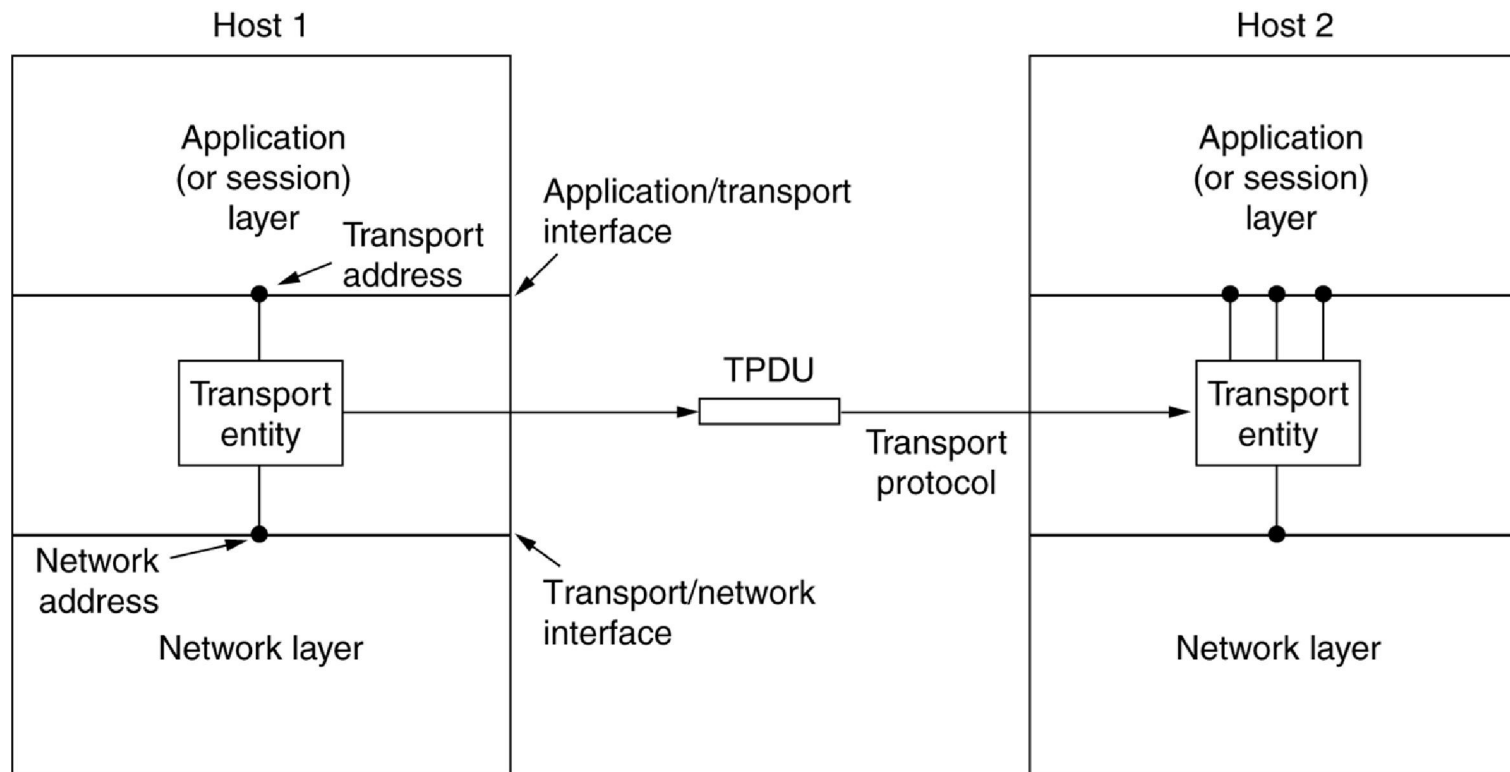
The Transport Entity

- The TCP software within the transport layer that implements the service will be referred to as the ***Transport Entity***:
 - This is the “Transport Service Provider”.
- The ***Transport Entity*** can be located in any number of places including:
 - Within the operating system (OS) kernel,
 - As a separate user process,
 - Within a library package bound to the network application,
 - On the network interface card.

The Transport Entity

- If the *protocol stack* is located within the OS the **primitives** are implemented as *system calls*:
 - These calls turn control of the machine over to the OS to send and receive the necessary PDUs.
- The next diagram shows the Transport Entity in context.
- As can be seen it shows the Transport Layer sitting between the Application and Network layers.

The Network, Transport and Application Layers



The Transport Entity

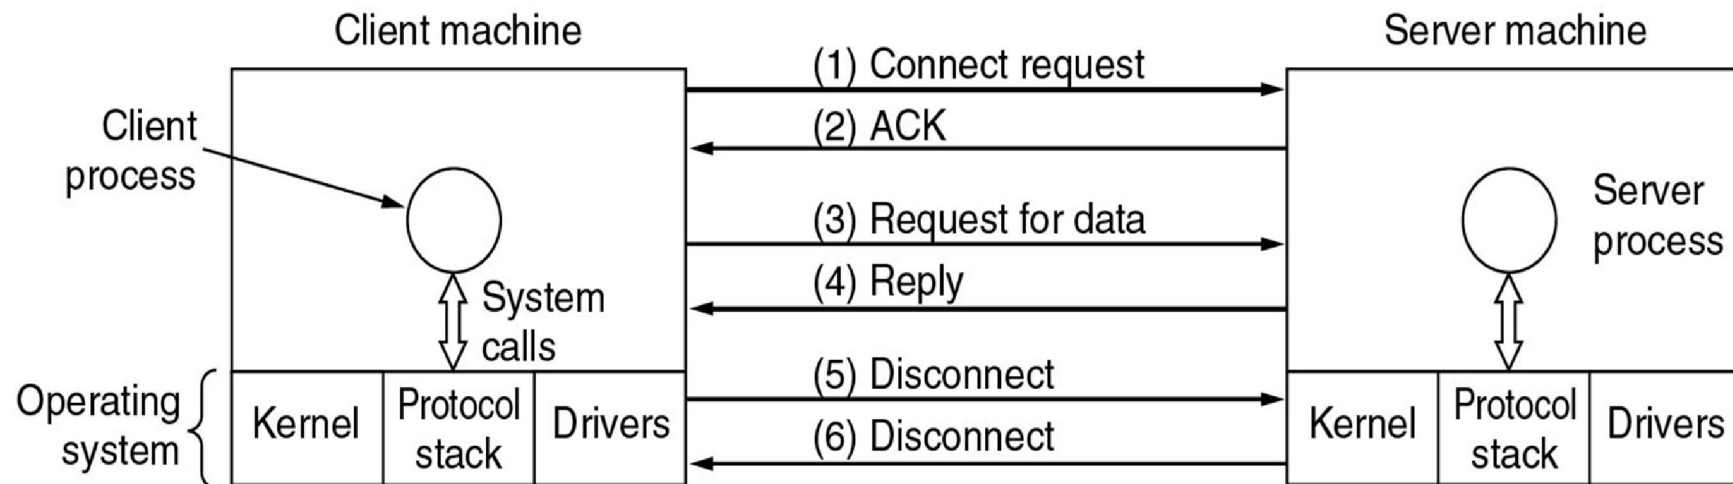
- It has a connection to each of the layers above and below:
 - It is the *Service Provider* to the Application Layer and,
 - A *Service User* of the Network Layer.
- The TPDU represents the framing structure that is exchanged between *Peer Entities*:
 - The PDU does NOT move horizontally between the Transport layers as depicted,
 - Instead it moves up-and-down through the Protocol Stack.

Transport Service Primitives

- What do these primitives look like?
- Consider a **generic** transport interface as shown in the next diagram:
 - Here can be seen a list of Primitives.
 - These primitives allow application programs to establish, use and release connections.
- Typically there is a very precise sequence of calls to these primitives:
 - The exact sequence differs for *clients* and *servers*.

A set of basic Transport Service primitives

Primitive	Packet sent	Meaning
LISTEN	(none)	Block until some process tries to connect
CONNECT	CONNECTION REQ.	Actively attempt to establish a connection
SEND	DATA	Send information
RECEIVE	(none)	Block until a DATA packet arrives
DISCONNECT	DISCONNECTION REQ.	This side wants to release the connection



Managing connections

- Before explaining the sequence of primitive calls it is important to understand the *Phases of Communication*.
- Recall from previous discussions on HDLC that there are generally three phases of communication associated with a connection-oriented service:
 - Phase 1: Connection Establishment,
 - Phase 2: Data Transfer, and,
 - Phase 3: Connection Release.
- During each phase a variety of PDUs are exchanged between the *client* and *server*:
 - Each PDU contains a message destined for the peer entity on the remote end of the channel.
 - The following slides explain the interaction for each phase.

Connection interactions per phase

- *Connection Establishment Phase:*
 1. The server application executes a LISTEN primitive (aka a ***call to Listen***).
 2. The server's *transport entity* responds to this primitive call by:
 - **Blocking** the server until a client request arrives.
 3. The client application then executes a CONNECT primitive (aka ***call to Connect***).
 4. The client's *transport entity* responds to this call by:
 - **Blocking** the client and sending a CONNECTION REQUEST TPDU to the Server.

Connection interactions per phase

- *Connection Establishment Phase:*
 5. Upon receipt of the CONNECTION REQUEST TPDU the server's *transport entity*:
 - Unblocks the server and returns a CONNECTION ACCEPTED TPDU to the Client,
 6. The client's *transport entity* then unblocks the client.
 7. The connection is now deemed **established**.

Connection interactions per phase

- *Data Transfer Phase:*
 - With an active connection now established data can be exchanged between the client and server using the SEND and RECEIVE primitives,
 - Each side must take turns using (blocking) RECEIVE and SEND.
- *Release Phase:*
 - Either the client or the server application can call a DISCONNECT primitive:
 - This causes a DISCONNECT TPDU to be sent to the remote transport entity.
 - Once the DISCONNECT TPDU has been received and acknowledged the connection is deemed **released**.

Berkeley Sockets

- The basic transport primitives discussed above are not standardised.
- Instead most OS designers have adopted the *socket primitives*:
 - These originated from the Berkeley University of California's UNIX OS which contained the original *TCP/IP* suite of internetworking protocols.
- Consequently the socket API has become the *de facto* standard for interfacing to TCP/IP

Origins of the Socket concept

- Coming from a UNIX background *sockets* use many concepts found in UNIX:
 - An application communicates through a *socket* in the same way that it transfers data to or from a file
- For File I/O UNIX uses an *open-read-write-close* paradigm:
 - An application makes the following calls in strict order:
 - *open* to prepare a file for reading/writing,
 - *read* or *write* to retrieve or send data to/from the file,
 - *close* to release the file.
- When *open* is first called a *descriptor* is returned:
 - All calls to the file use this *descriptor*,
 - Socket communication also uses this ***descriptor*** approach.

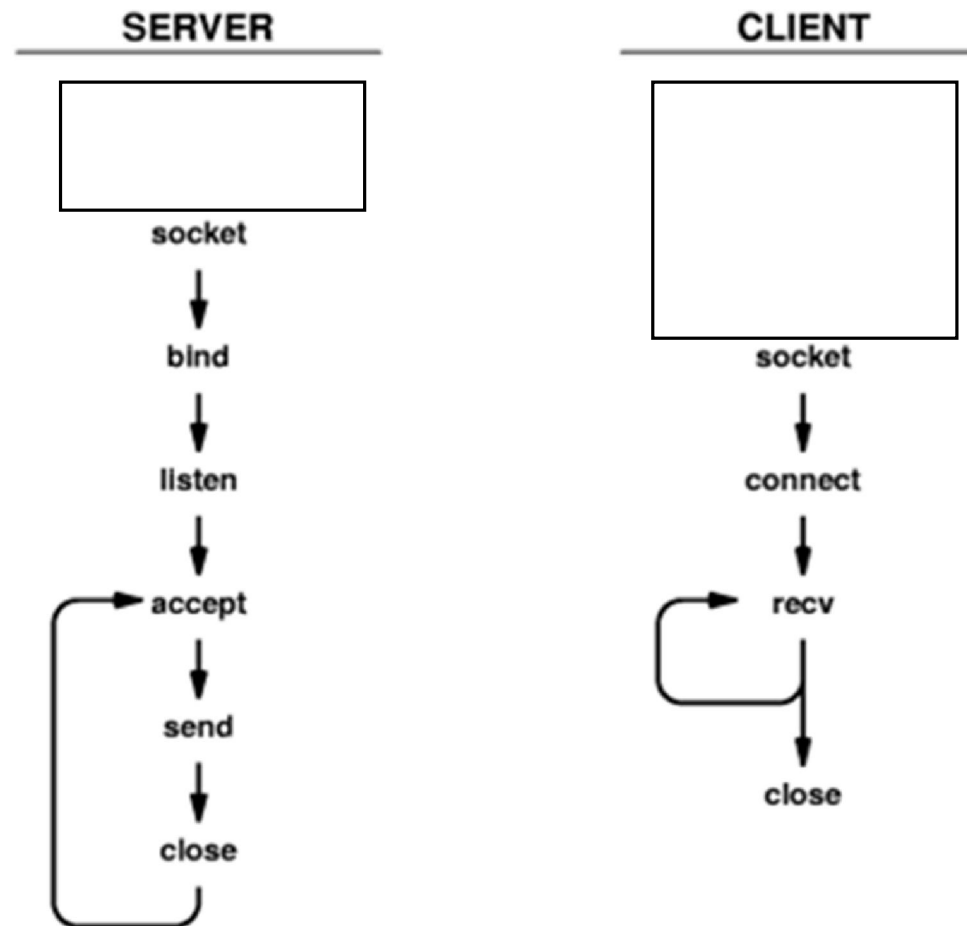
Socket Communication in UNIX

- Applications that need to use the TCP/IP protocols to communicate must request the OS to create a **socket**:
 - This is an abstract concept which will be explained in detail later.
- Similar to File I/O the OS returns a ***descriptor*** that uniquely *identifies* the socket:
 - As with File I/O this descriptor must be used in all interactions with the socket.
- The following slides lists the Socket API primitives and illustrates an example of how these primitives are used by a client and server application

The Socket Transport Primitives

Primitive	Meaning
SOCKET	Create a new communication end point
BIND	Attach a local address to a socket
LISTEN	Announce willingness to accept connections; give queue size
ACCEPT	Block the caller until a connection attempt arrives
CONNECT	Actively attempt to establish a connection
WRITE/SEND	Send some data over the connection
READ/RECEIVE	Receive some data from the connection
CLOSE	Release the connection

An Example Client-Server Interaction Using Sockets



The Socket Primitives - explained

- The following primitives are executed by *servers*:
 - SOCKET: This primitive creates a new *end point* within the Transport Entity:
 - Table space is allocated within the transport entity,
 - A file descriptor is returned which is used in all future calls
 - BIND: This primitive binds a socket to a network address:
 - This allows remote clients to connect to it

The Socket Primitives - explained

- LISTEN: This primitive allocates a queuing space within the transport entity for incoming call requests.
- ACCEPT: This primitive blocks the server waiting for an incoming connection:
 - Upon receipt of a connection request the *transport entity* creates a new socket identical to the original one and returns a file descriptor to the server.
 - The server *forks off* a new process or *service thread* to handle the connection on the new socket
 - The server also continues to wait for more connections on the original socket

The Socket Primitives - explained

- The following primitives are executed by *clients*:
 - SOCKET: As before
 - CONNECT: This primitive blocks the client and actively starts the connection process
- The following primitives are executed by *clients* and *servers*:
 - SEND and RECV: These primitives are used to transmit and receive data over the full-duplex connection
 - CLOSE: This primitive releases the transport connection

Sockets and Socket Libraries

- In most systems the *socket* functions are part of the OS.
- Some systems, however, require a *socket library* to provide the interface to the *transport entity*:
 - These operate differently to a *native* socket API,
 - The code for the library socket procedures are linked into the application program and resides in its address space,
 - Calls to a socket library pass control to the library routine as opposed to the OS.
- Both implementations provide the same semantics from a programmer's perspective
- Applications using either implementation can be **ported** to other computer systems.

Example use of Sockets

- The next slide shows an example Client application using the Sockets API.
- It is a simple Daytime Client which connects to a Daytime Server for the purpose of retrieving the current date and time from the Server Host.
- This application is written in 'C'.
- Lines 24, 41, 44 and 57 show four of the socket primitives being called.
- This programme will be explained in class and you will have a chance to compile and run it against a pre-compiled server.

An Example Client Using Sockets

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <string.h>
4  #include <unistd.h>
5  #include <sys/types.h>
6  #include <sys/socket.h>
7  #include <netinet/in.h>
8  #include <arpa/inet.h>
9  #include "Practical.h"
10
11  int main(int argc, char *argv[]) {
12      char recvbuffer[BUFSIZE]; // I/O buffer
13      int numBytes = 0;
14
15      if (argc < 3) // Test for correct number of arguments
16          DieWithUserMessage("Parameter(s)",
17                              "<Server Address> <Server Port>");
18
19      char *servIP = argv[1]; // First arg: server IP address (dotted quad)
20
21      in_port_t servPort = atoi(argv[2]);
22
23      // Create a reliable, stream socket using TCP
24      int sock = socket(AF_INET, SOCK_STREAM, IPPROTO_TCP);
25      if (sock < 0)
26          DieWithSystemMessage("socket() failed");
27
28      // Construct the server address structure
29      struct sockaddr_in servAddr; // Server address
30      memset(&servAddr, 0, sizeof(servAddr)); // Zero out structure
31      servAddr.sin_family = AF_INET; // IPv4 address family
32      // Convert address
33      int rtnVal = inet_pton(AF_INET, servIP, &servAddr.sin_addr.s_addr);
```


An Example Client Using Sockets

```
34     if (rtnVal == 0)
35         DieWithUserMessage("inet_pton() failed", "invalid address string");
36     else if (rtnVal < 0)
37         DieWithSystemMessage("inet_pton() failed");
38     servAddr.sin_port = htons(servPort);    // Server port
39
40     // Establish the connection to the echo server
41     if (connect(sock, (struct sockaddr *) &servAddr, sizeof(servAddr)) < 0)
42         DieWithSystemMessage("connect() failed");
43
44     while ((numBytes = recv(sock, recvbuffer, BUFSIZE - 1, 0)) > 0) {
45         recvbuffer[numBytes] = '\0';    // Terminate the string!
46         fputs(recvbuffer, stdout);    // Print the echo buffer
47         /* Receive up to the buffer size (minus 1 to leave space for
48          * a null terminator) bytes from the sender */
49     }
50     if (numBytes < 0)
51         DieWithSystemMessage("recv() failed");
52     //     else if (numBytes == 0)
53     //         DieWithUserMessage("recv()", "connection closed prematurely");
54
55     fputc('\n', stdout); // Print a final linefeed
56
57     close(sock);
58     exit(0);
59 }
60
```