

IP Addresses and Hostnames

- The applications developed thus far use *IP addresses*:
 - Typically an IP address is passed from the command line to be used by a client application to establish a connection to a remote server,
 - Practically, however, end-users rarely know the IP address of the remote server applications,
 - Instead, end-users use human-readable ***hostnames***.
- Advantages for using *hostnames*:
 - Using *hostnames* instead of IP addresses is easier for end-users,
 - A host's IP address can change without affecting its *hostname*
- However, client applications connecting to remote server applications still require an IP address:
 - Some method is required to map hostnames to IP addresses.

Hostname-to-IP Address Mappings

- A *Name Service* is used to map between *hostnames* and *IP addresses* (amongst other things):
 - The process of mapping a *hostname* to a numeric quantity such as, an IP address or, Port Number, is called ***resolution***,
 - When an IP address for a particular hostname is obtained from a *name service*, the hostname is said to be ***resolved***.
- Two primary name service sources are:
 - The *Domain Name System* (DNS). This is a distributed name service requiring the use of the DNS protocol. and,
 - Local configuration databases which are operating-system specific.

Hostname-to-IP Address Mappings

- Fortunately from a programming perspective the details of the name service are hidden:
 - Programmers only need to know how to ask for a name to be **resolved**.

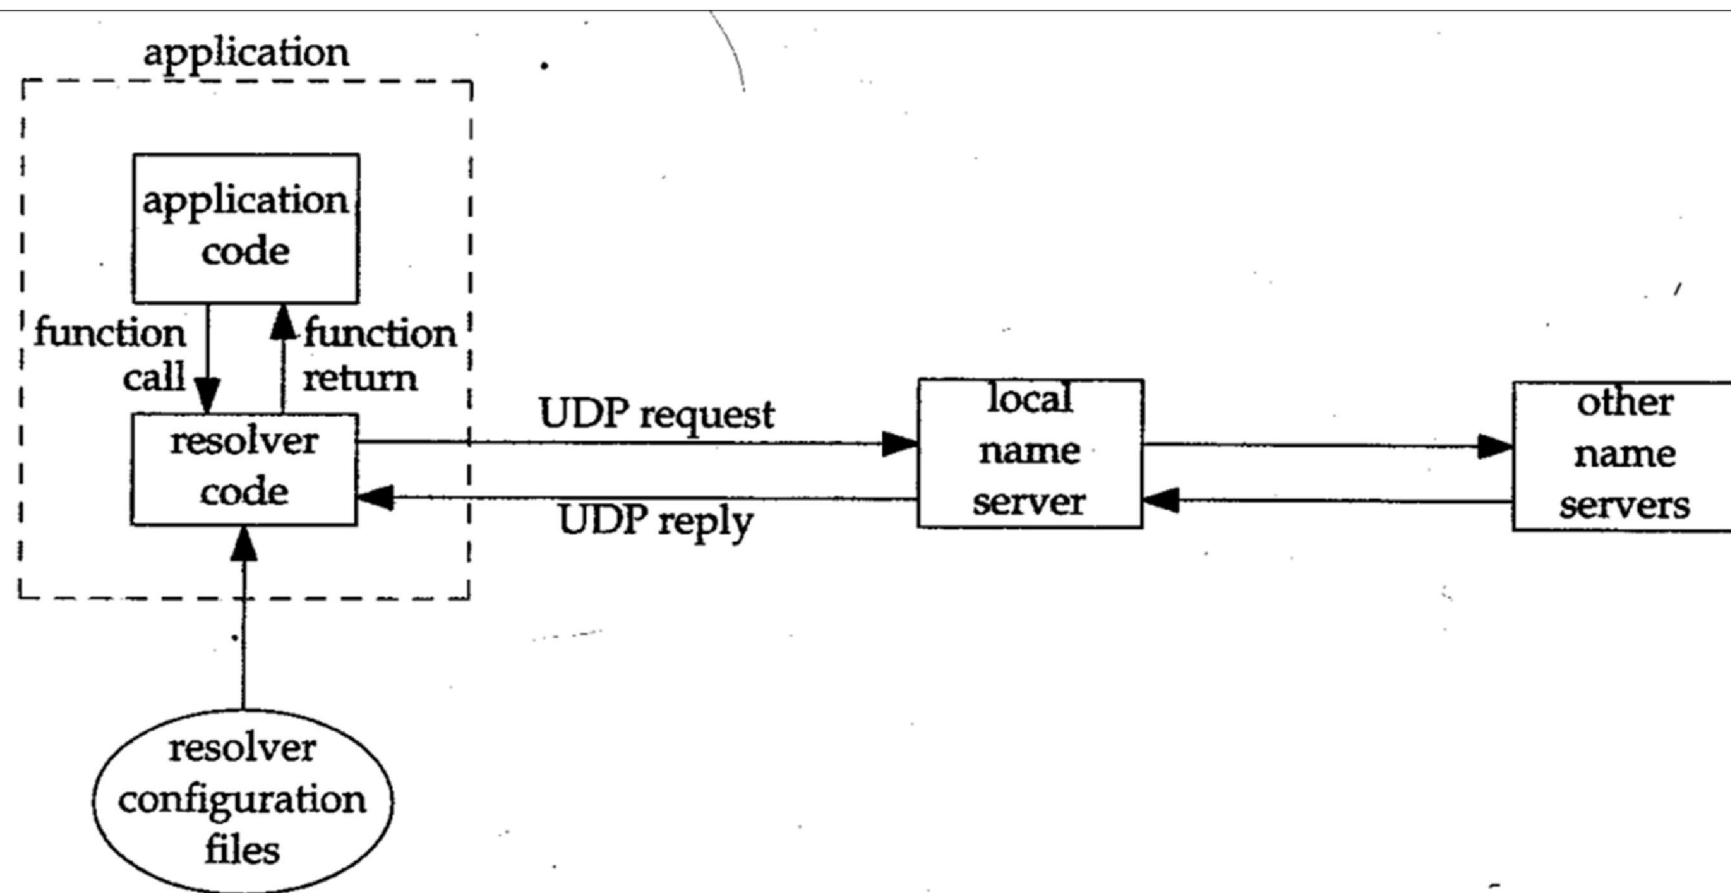
Resolvers and Name Servers

- Organizations often run one or more *name servers* a.k.a. DNS (Domain Name System) servers.
- Applications interact with DNS servers using functions imported from a library known as a ***resolver***:
 - The *resolver* code is contained in a system library and is link-edited into the application during the *build* process,
- Calls to the *resolver* code are made using functions such as ***getaddrinfo()*** and ***getnameinfo()***
 - The former maps a *hostname* into its IP address, and the latter does the reverse mapping

Operation of Resolvers

- Prior to contacting a name server, the resolver code refers to a local *configuration file* to determine the IP Address of the name server(s):
 - The file `/etc/resolv.conf` normally contains the IP address of the local name servers
- The resolver then sends a query to a local name server for a *resource record*:
 - Entries in the DNS are known as *resource records* (RRs)
 - If necessary, the local name server may query another name server for the RR.

Resolvers and Name Servers



Name and Address Conversions

- There are a number of different *resource records*:
 - *A* records map hostnames to a 32-bit IPv4 address,
 - *AAAA* records map hostnames to a 128-bit IPv6 address,
 - *PTR* records map *IP* addresses into *hostnames*,
 - *MX* records specifies a host to act as a *mail exchanger*,
 - *CNAME* records map common services, such as *ftp* and *www* to the actual host providing the service
 - Example www.dit.ie has the *canonical* name *remus.dit.ie*.
- The RRs that we are interested in is the ***A record***.

The *getaddrinfo()* function

- The ***getaddrinfo()*** function performs a query for an *A* record. It is defined as follows:

```
int getaddrinfo (const char *hostStr, const char *serviceStr,  
                 const struct addrinfo *hints, struct addrinfo **results);
```

e.g. **getaddrinfo** ("www.google.com", 0, NULL, &addrList);

Returns: NULL if OK and a non-null error code if unsuccessful.

- The function returns one, or more, *addrinfo* structures, each of which contains an Internet address that can be specified in subsequent calls to *bind()* or *connect()*.

getaddrinfo() parameters

- ***hostStr:***
 - Points to a null-terminated character string representing a **host name** such as *aisling*, or, a fully qualified domain name (FQDN) such as: *aisling.student.dit.ie*
- ***serviceStr:***
 - A service name which is translated to the corresponding port number. This is ignored for the moment
- ***hints:***
 - filters to be used to restrict the information returned
- ***results:***
 - the address of a pointer (type *struct addrinfo*) to hold the first address of a linked list of results i.e. the protocol addresses

addrinfo structure

- Each entry in the linked list is an *addrinfo* structure which is defined as follows:

```
struct addrinfo {  
    int          ai_flags;    // Flags to control information resolution  
    int          ai_family;  // Family: AF_INET, AF_UNSPEC etc.  
    int          ai_socktype; // Socket type: SOCK_STREAM,  
                                SOCK_DGRAM  
    int          ai_protocol; // Protocol: 0 (default) or IPPROTO_XXX  
    socklen_t    ai_addrlen; // Length of socket address ai_addr  
    struct sockaddr *ai_addr; // Socket address for socket  
    char         *ai_canonname; // Canonical name  
    struct addrinfo *ai_next; // Next addrinfo in linked list  
};
```

addrinfo members of interest

- There are five members of interest:
 - **ai_family**: Specifies the address family supported by the hostname. Recall we are interested in AF_INET which is the TCP/IP stack.
 - **ai_socktype**: Specifies the type of socket supported by the hostname. Recall we are interested in SOCK_STREAM i.e. a streaming socket.
 - **ai_protocol**: Specifies the specific protocol supported by the hostname. Recall we are interested in IPPROTO_TCP i.e. TCP
 - ***ai_addr**: Points to a structure that holds the full TCP/IP address for the hostname. The next slide refers to this structure.
 - **ai_addrlen**: The length of the socket address.

ai_addr and sockaddr_in

- ***ai_addr** points to a socket address structure of type *sockaddr*. This is the **generic** address structure.
 - TCP applications use the **family-specific** address structure (*sockaddr_in* – members shown below) which is then typecast to the generic address structure in any calls to the socket primitives.

```
struct sockaddr_in {  
    uint8_t        sin_len;           // length of structure (16)  
    sa_family_t     sin_family;      // AF_INET  
    in-port_t       sin_port;        // 16-bit TCP or UDP port number  
                                     // network byte ordered  
    struct in_addr   sin_addr;        // 32-bit IPv4 address  
                                     // network byte ordered  
    char            sin_zero[8];      // unused  
};
```

ai_addr and *sockaddr_in*

- Fortunately, it is not necessary to access individual members of the generic socket address structure:
 - Calls to socket primitives that need addressing information only require a pointer to the structure.
 - For instance a typical call to `connect()` is as follows:
`connect(sock, (struct sockaddr *) &servAddr, sizeof(servAddr)),`
 - Where *servAddr* is a struct of type *sockaddr_in*
- To use information returned by ***getaddrinfo()*** we can use:
`connect(sock, addr->ai_addr, addr->ai_addrlen),`
 - Where *addr* is a pointer to a structure of type *addrinfo* and, *ai_addr* (a member of *addr*) is a pointer to a generic address structure.

getaddrinfo() and *freeaddrinfo()*

- Eventually the linked list of results returned by ***getaddrinfo()*** must be deallocated:
 - This requires the use of the auxiliary function, ***freeaddrinfo()***
 - Given a pointer to the head of the linked list, ***freeaddrinfo()*** frees all the storage allocated for the list.
 - Failure to call this function can result in a memory leak
- The following example program (GetAddrInfo.c) illustrates the use of ***getaddrinfo()*** and ***freeaddrinfo()***:
 - The program takes two command-line parameters, a hostname and a service name (or port number), and prints the IP address(es)
./GetAddrInfo www.tudublin.ie http

A sample program using *getaddrinfo()*

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <string.h>
4  #include <netdb.h>
5  #include "Practical.h"
6
7  int main(int argc, char *argv[]) {
8
9      if (argc != 3) // Test for correct number of arguments
10         DieWithUserMessage("Parameter(s)", "<Address/Name> <Port/Service>");
11
12     char *addrString = argv[1]; // Server address/name
13     char *portString = argv[2]; // Server port/service
14
15     // Tell the system what kind(s) of address info we want
16     struct addrinfo addrCriteria; // Criteria for address match
17     memset(&addrCriteria, 0, sizeof(addrCriteria)); // Zero out structure
18     addrCriteria.ai_family = AF_INET; // The TCP/IP address family
19     addrCriteria.ai_socktype = SOCK_STREAM; // Only stream sockets
20     addrCriteria.ai_protocol = IPPROTO_TCP; // Only TCP protocol
21
22     // Get address(es) associated with the specified name/service
23     struct addrinfo *addrList; // Holder for list of addresses returned
24     // Modify servAddr contents to reference linked list of addresses
25     int rtnVal = getaddrinfo(addrString, portString, &addrCriteria, &addrList);
26     if (rtnVal != 0)
27         DieWithUserMessage("getaddrinfo() failed", gai_strerror(rtnVal));
28
29     // Display returned addresses
30     for (struct addrinfo *addr = addrList; addr != NULL; addr = addr->ai_next) {
31         PrintSocketAddress(addr->ai_addr, stdout);
32         fputc('\n', stdout);
33     }
34
35     freeaddrinfo(addrList); // Free addrinfo allocated in getaddrinfo()
36
37     exit(0);
38 }
```

Code explained

- Line 16 is a *struct* that is used to restrict the type of information to be returned:
 - In this case we are only interested in TCP/IP addresses
 - The address of this *struct* is passed as the third argument (**hints**) to **getaddrinfo()**
- Line 25 is the call to **getaddrinfo()**
- Lines 30 iterates over each node of the linked list
- Line 31 prints the IP address and Port number from the *ai_addr* member of the current linked list node:
 - **ai_addr** points to a *struct* of type **sockaddr**, the generic address structure.
 - The addresses are then taken from the **sa_data** member.
 - Recall we used a structure of type **sockaddr_in** which has members: **sin_addr** and **sin_port**