

Racap of Flow Control in the Data Link layer

- Recall the use of the Sliding Windows Flow Control technique:
 - This technique allows *multiple* **Frames** to be in transit in sucession,
 - This provides for more *efficient* **Link Utilization**.
- Characteristics of the technique are:
 - Both stations use an extended buffer size to hold *multiple* frames.
 - The Sending/Receiving stations maintain a list of frames already sent/received.
- The transmission link is effectively treated as a *pipeline* that can be *filled* with many frames in transit.


Example Sliding Windows Flow Control

- Consider two Stations, A and B exchanging data. Assume Station A is sending data to Station B:
 - Station A is the **Sender** and Station B is the **Receiver**.
- Before any data is exchanged, each station allocates buffer space for W frames (for example consider **$W=8$**)
 - This means that Station B can *accept* up to **W** (8) frames and, Station A can *send* up to **W** (8) frames, without individual frame acknowledgements (ACKs) being sent or received.

Example Sliding Windows Flow Control

- All frames contain a *sequence number*:
 - All frames from Station A to Station B contain a sequence number for the **current frame**,
 - All *acknowledgements* from Station B contain the sequence number of the **next frame** expected.
- Frames leaving Station A are stored in an **outgoing** buffer on Station A until an ACK is **received**.
- Frames arriving at Station B are stored in an **incoming** buffer on Station B until an ACK is **sent**.

Example Sliding Windows




Sliding Windows Flow Control


- *Multiple* frames can be *acknowledged* using a single control message (*implicit ACK*):
 - e.g. Receipt of ACK for frame **2** (RR3) followed later by ACK for frame **5** (RR6) *implies* acknowledgement of frames **3** and **4**.
- Station A maintains a list of frame numbers it is allowed to *send* and, Station B maintains a list of frame numbers it is prepared to *receive*:
 - Each list cannot extend beyond the window size.
 - These lists can be considered as *windows*.

Sliding Windows Flow Control

Transmitter's Perspective



Receiver's Perspective



Error Control

- Recall the purpose of Error Control:
 - Sender and Receiver stations co-ordinate activities to recover from **Lost** or **Damaged** Frames etc.
- Error Control involves enhancing Flow Control techniques with additional functionality such as:
 - **Transmission Timers.** Sender stations set a timer for each frame transmitted and takes action when a timer expires.
 - **Negative ACKs.** A Receiver station can reject an *out-of-sequence/damaged* frame with a REJ(5) or SREJ(4) message.
 - Example Error Control techniques include: **Go-Back-N** and **Selective Reject**:
 - Both techniques are based on the Sliding Windows Flow Control technique.

TCP Error/Flow Control


- *TCP* also uses an Error Control technique based on **Sliding Window Flow Control** technique.
- It is different to that used in the Data Link layer as follows:
 - Data is sent in **Segments** not **Frames**.
 - Sequence numbers relate to **Bytes** not Segments. Each **byte** in a segment is numbered:
 - Each Segment identifies the first byte in the data field.
 - ACKs contain the number of the next byte expected.
 - There are no **Negative** ACKs.

TCP Error/Flow Control

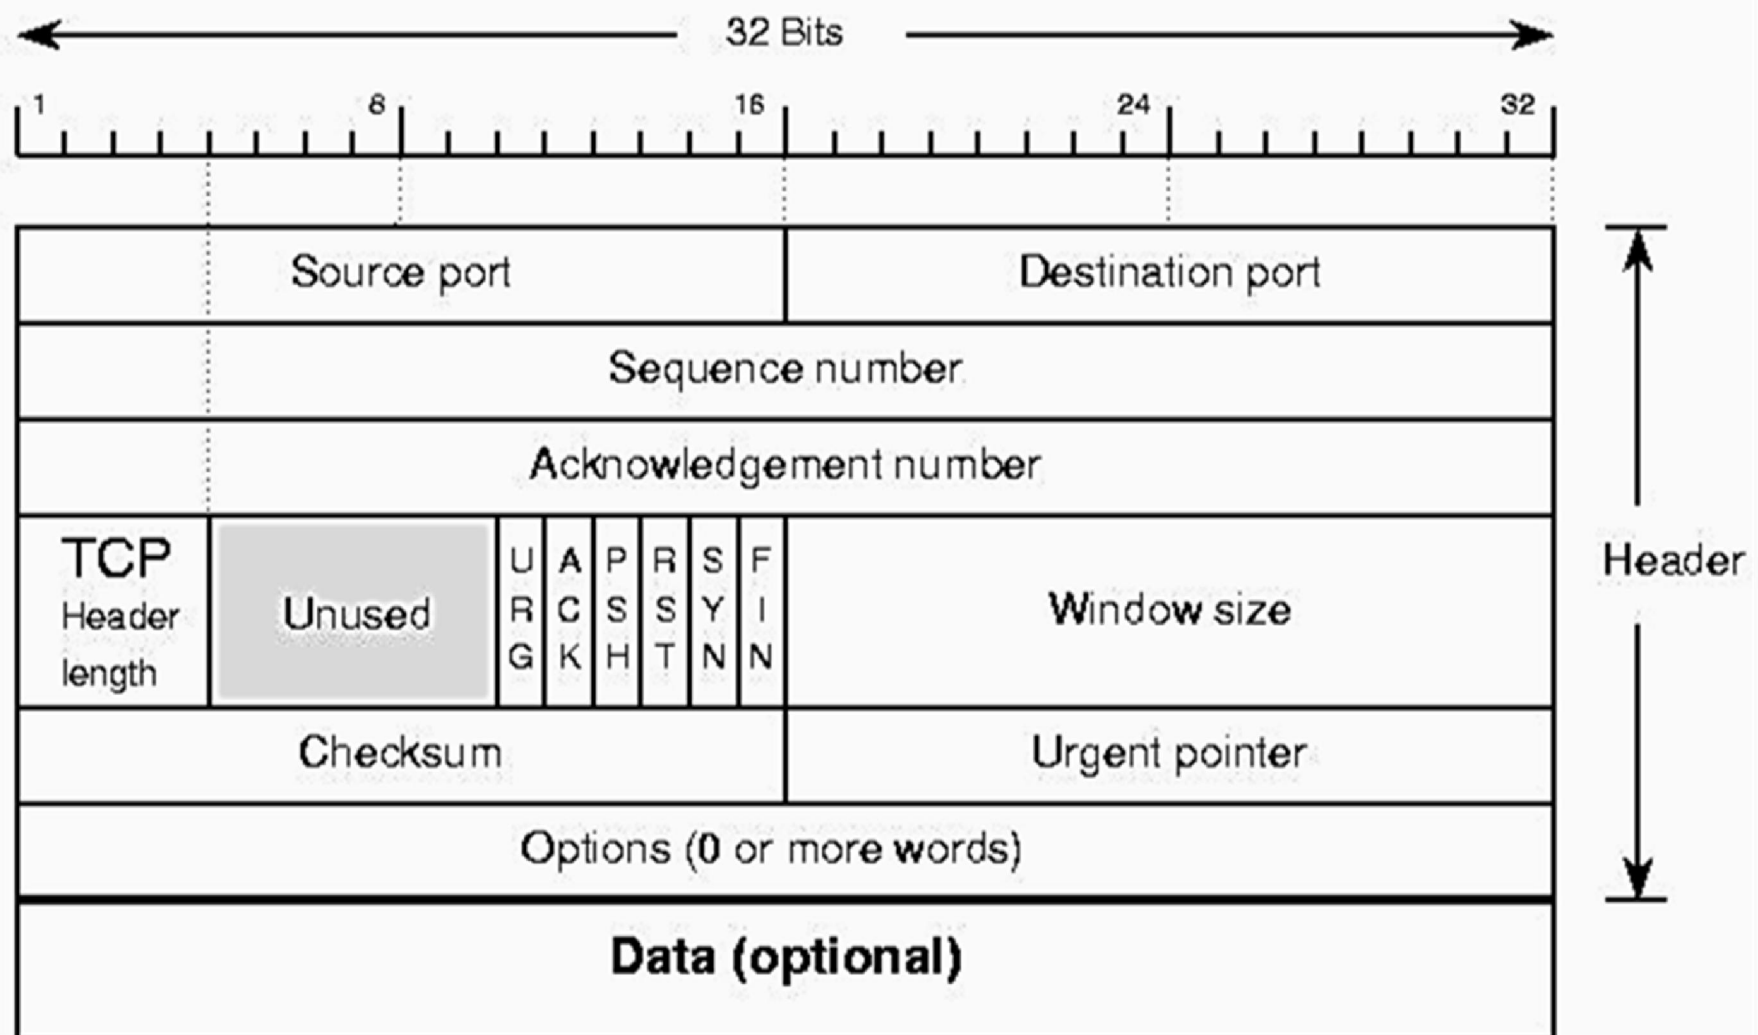
- Senders and Receivers maintain lists of bytes sent/received.
- Buffers used to hold incoming segments are measured in bytes (not segments).

TCP Sliding Windows

- Here it can be seen that the available buffer space *decreases* as data is stored in the buffer and *increases* as data leaves the buffer:



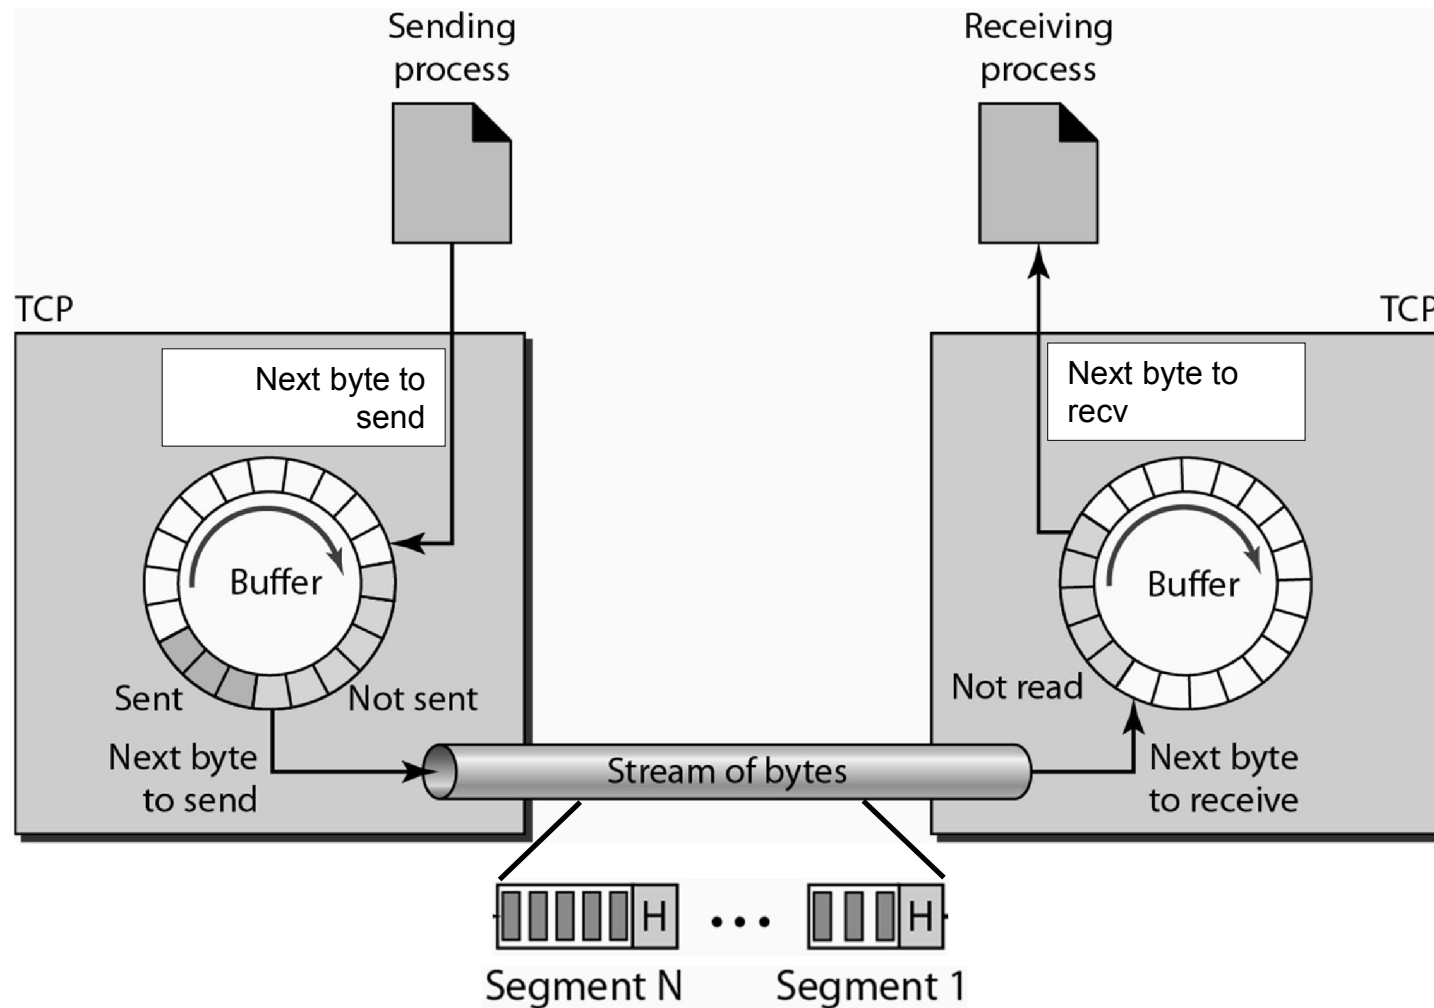
TCP Segment Format



TCP Flow Control – *Buffers and Windows*

- Recall that TCP creates two buffers per socket:
- These can be viewed with the **netstat** utility:
 - One for incoming data (known as **RECV-Q** in netstat)
 - One for outgoing data (known as **SEND-Q** in netstat)
- Incoming buffers can easily overflow.
- To prevent this, the receiving TCP entity uses a ***Window Mechanism***.

TCP Internal Data Buffers



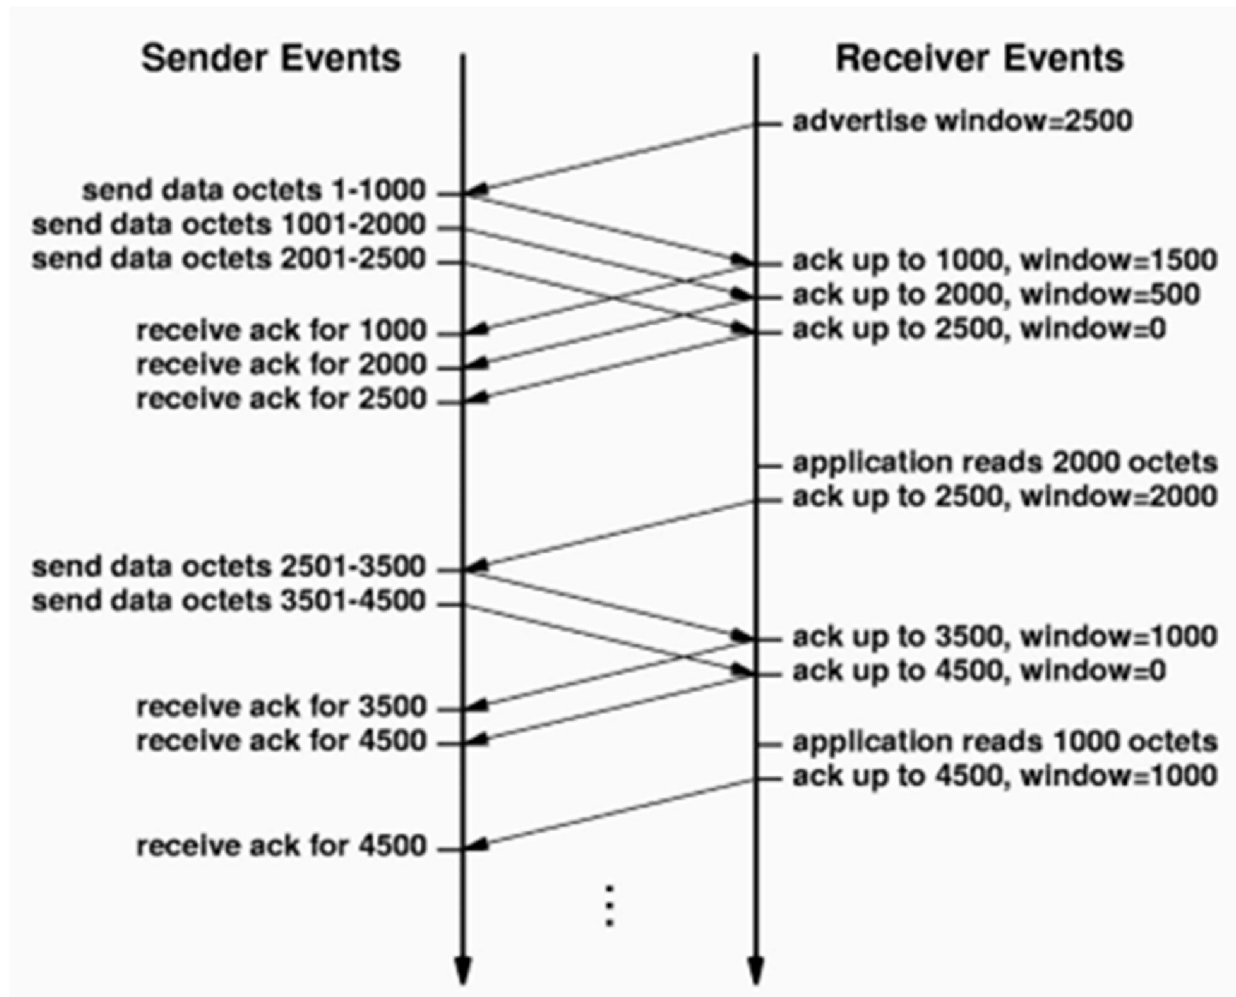
TCP Flow Control – *Buffers and Windows*

- Each end of the connection allocates a ***window*** (**RECVQ** buffer) to hold incoming data:
 - The size of the initial window, in bytes, is set using the *Window Size* field during *Phase 1 (Connection Initialization)* when both sides exchange SYN messages.
 - This is known as a **Window Advertisement**.

TCP Flow Control – *Buffers and Windows*

- Thereafter, throughout *Phase 2 (Data Exchange)*, all **ACKs** messages include a ***Window Advertisement***:
 - Again using the *Window Size* field.
 - The *window advertisement* can be positive or zero depending on the available space in **RECV-Q**.

Operation of *Window Advertisements*



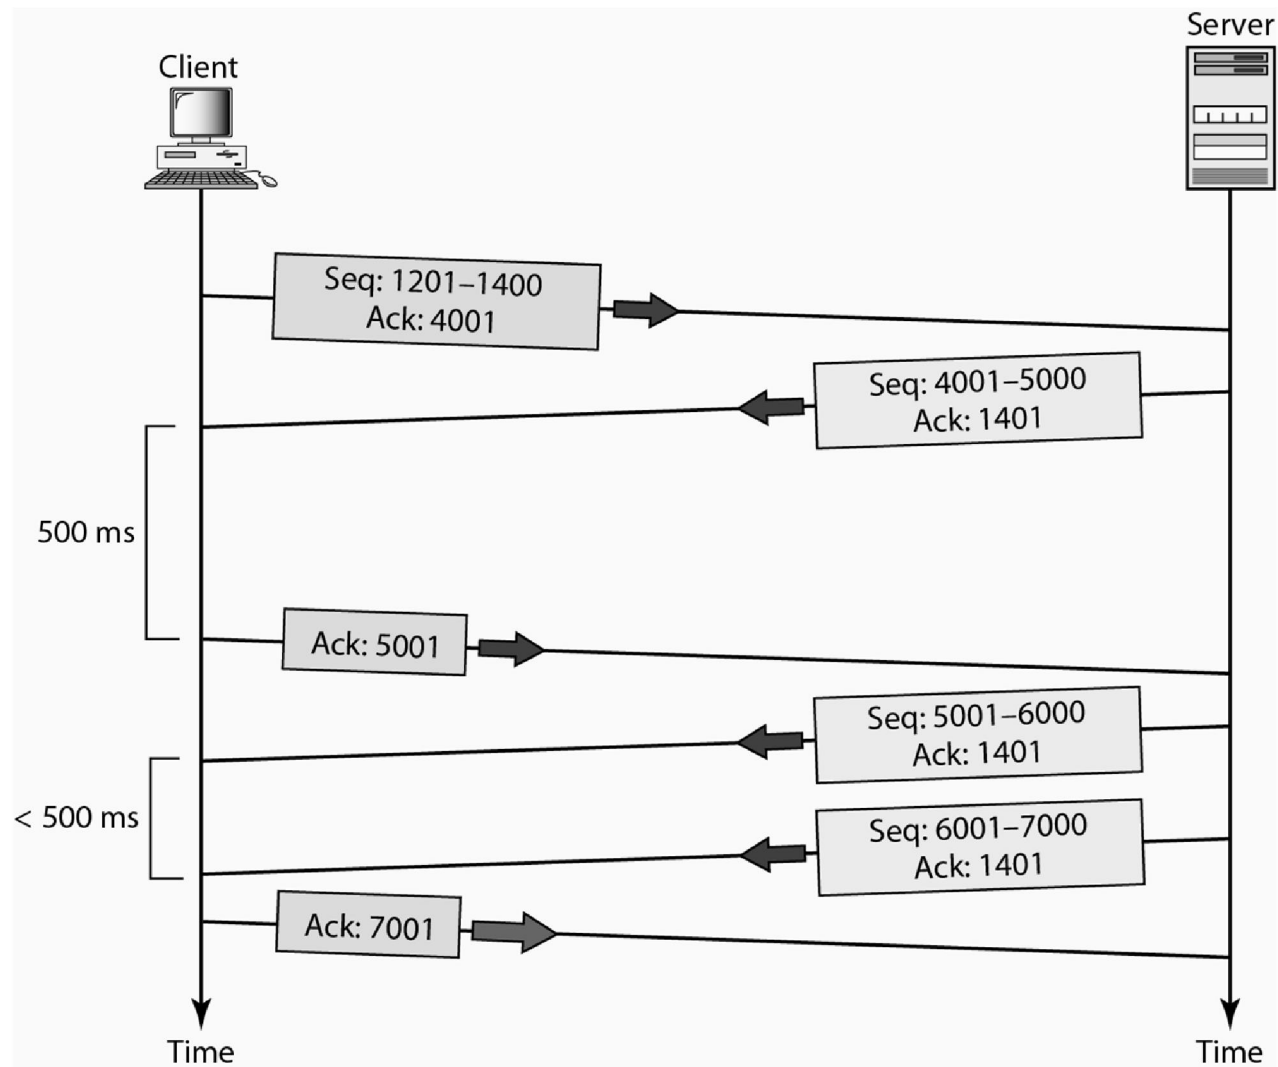
Achieving Reliability

- In addition to Flow Control, TCP must address the following reliability problems:
 - Unreliable delivery by the underlying communication system:
 - Segments can be **lost**, **duplicated**, **delayed**, or delivered **out-of-order** by the underlying communication system (IP Layer).
 - Computer reboot:
 - During an active connection, either side may re-boot unexpectedly.
 - Segments arriving after re-boot need to be dealt with.

TCP Error Control

- This requires the use of some form of ***error control***.
- Interestingly to implement *flow* and *error control*, TCP is only equipped with two elements:
 - **Positive ACKs** and **Timers**,
 - Significantly TCP does not have a **Negative ACK**.
- The following slide shows the operation of TCP during normal *Data Exchange* without error.

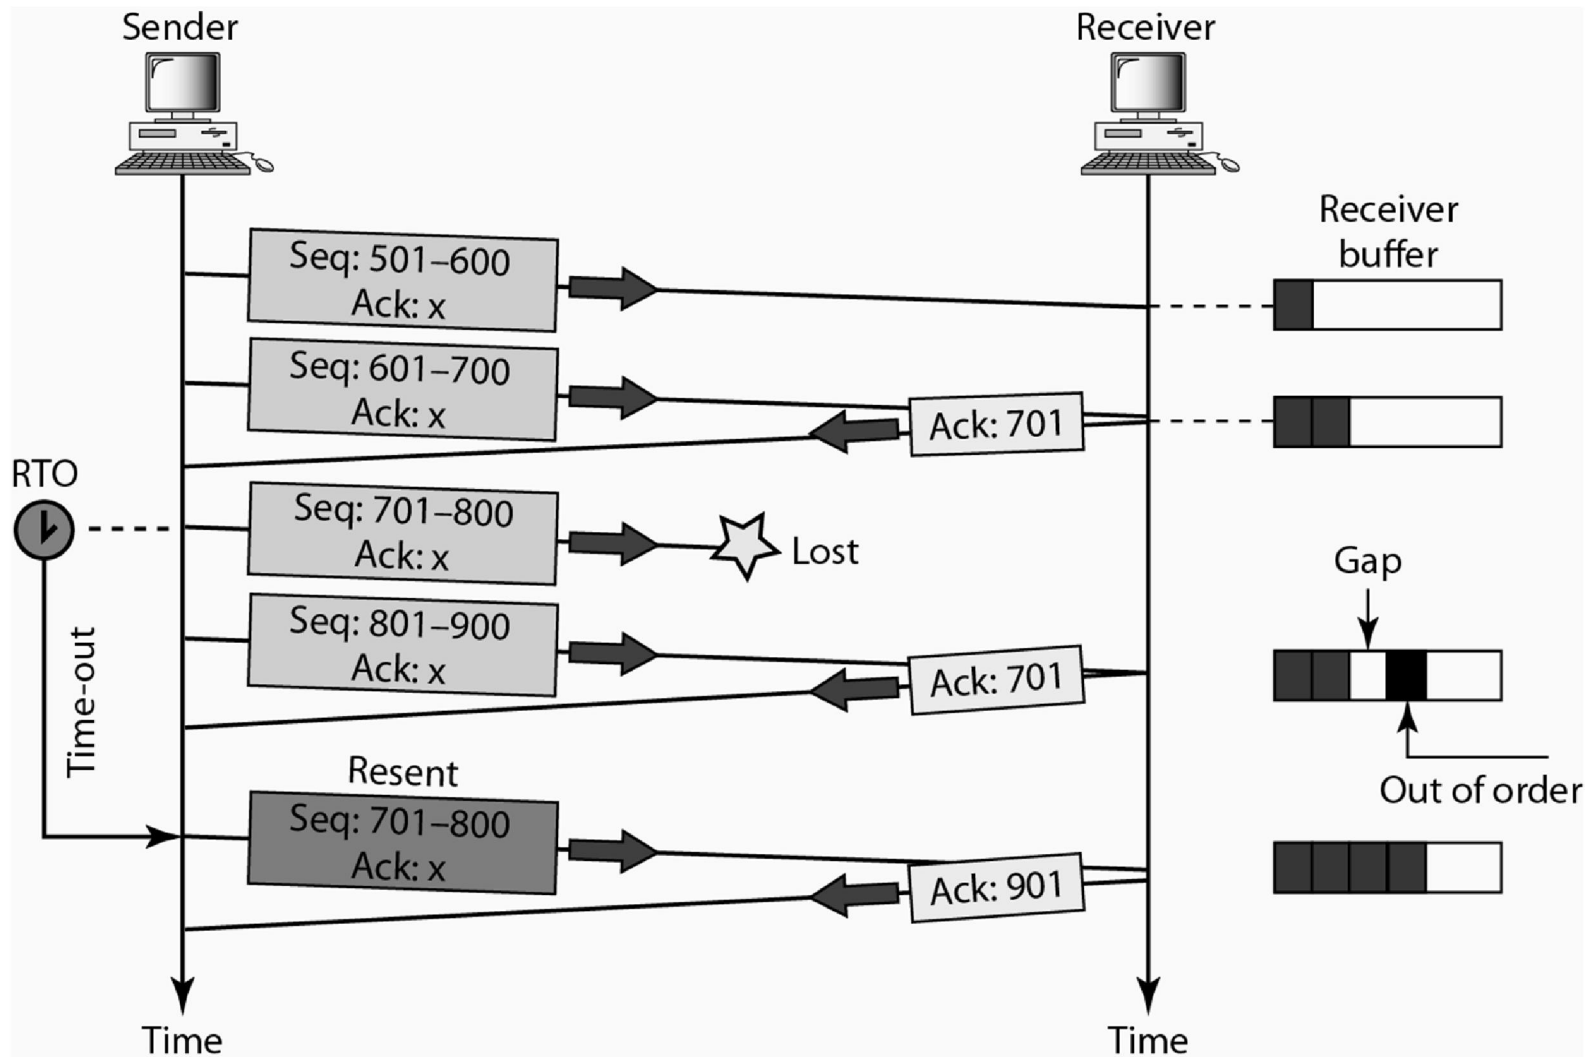
TCP Data Flow – Normal Operation



TCP Error Control - Segment Loss

- To deal with *lost* or *out-of-sequence* segments, TCP implements a **Retransmission** scheme:
 - This involves the retransmission of ACKs and/or lost segments.
- For the Sender, **timers** play a key role in error control:
 - Upon expiry of a timer (relating to an unacknowledged segment) the Sender simply re-transmits the segment,
 - A key question is “How long should TCP wait before retransmitting?”
- For the Receiver, sequence numbers play a key role in detecting **lost** or **out-of-sequence** segments:
 - A previous ACK is re-transmitted in response until the situation is rectified.

TCP Data Flow – Lost Segment Scenario



Factors affecting TCP's Retransmission scheme

- How long should a Sending TCP entity wait before re-transmitting a segment?
- The answer depends upon two factors:
 - The underlying network architecture, and,
 - Traffic levels across the network.
- TCP takes a measure of the delay between sending a segment and receiving an ACK.
 - This is known as Round-Trip Time (**RTT**).
- TCP uses the value for RTT to determine an appropriate value for the re-transmission timer (**RTO** - Retransmission Timeout).

Calculation of Retransmission Timeout (RTO)

- For each active connection, the RTT is continuously monitored:
 - It represents the *network latency*.
- TCP **adapts** its RTO timer to match the varying RTT values across the network:
 - It uses a weighted average of RTT and a variance factor,
 - It continuously re-calculates a value for RTO.

TCP's *Adaptive Retransmission* Scheme

- This is known as an *adaptive retransmission* scheme and is the key to TCP's success.
- This adaptability helps TCP to react quickly to changes in *traffic levels* and to maximize throughput on each connection.